

Web Data Management Development Project

Xiaoxu Gao

Alexander Overvoorde

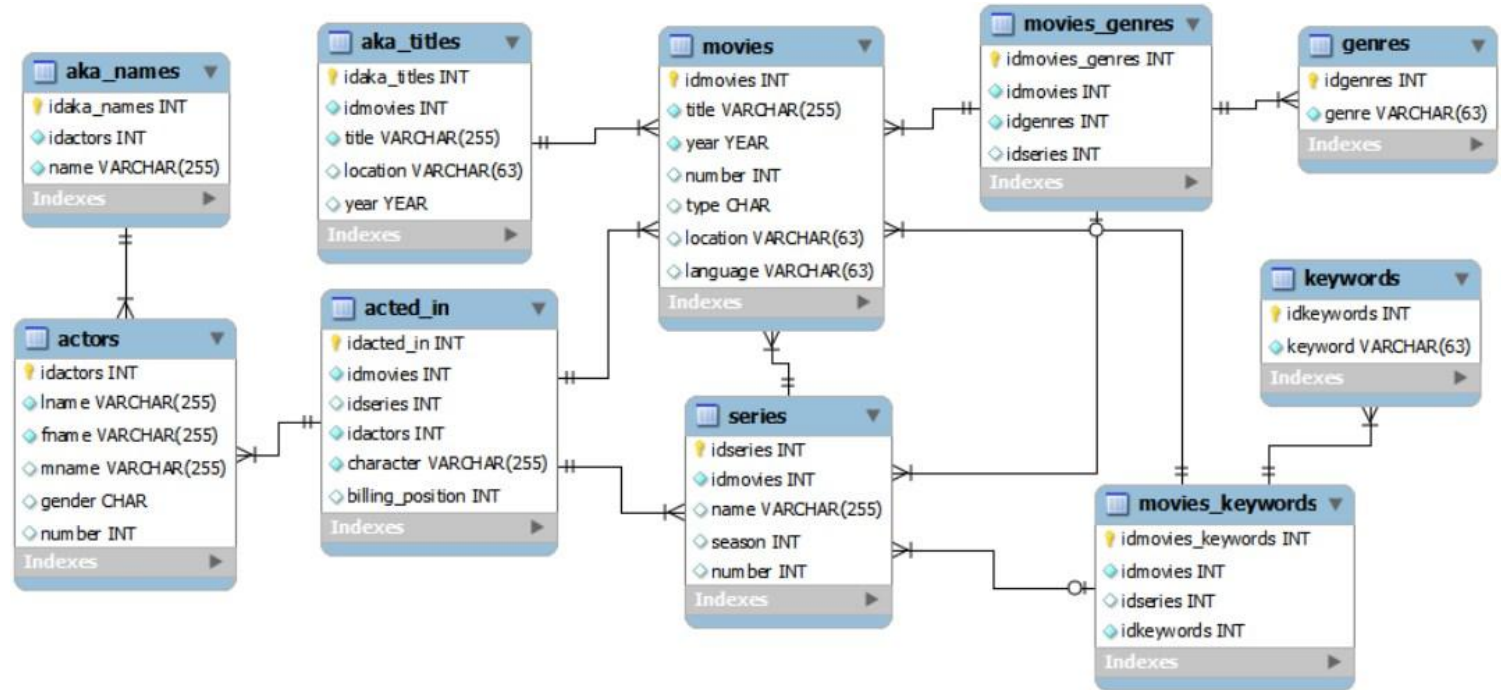
Tom Peeters

Piotr Tekieli

Choice of technologies

- PostgreSQL
- MongoDB (document-based)
- Neo4j (graph-based)

PostgreSQL - Schema



PostgreSQL - Query design 1/5

- For each SC, we worked around a different “central point” :
- SC1 - Movies
- SC2-3 - Actors
- SC4-5 - Genres
- Further connections with a related information were obtained with JOINS.

PostgreSQL - Query design 2/5

- Each JOIN is written as separate SQL query, therefore, to get <Full Movie Info> we need to send 5 queries to the DB.

Example : (based on movie id) Get Actor Info + Character Played + His Billing Position

```
SELECT DISTINCT a.idactors, a.fname, a.lname, a.gender, ai.character, ai.billing_position  
FROM actors a  
JOIN acted_in ai ON a.idactors = ai.idactors  
JOIN movies m ON m.idmovies = ai.idmovies  
WHERE m.idmovies = <someid> ORDER BY ai.billing_position;
```

PostgreSQL - Query design 3/5

- We used built-in functions (i.e. COUNT) to create statistics for genres and actors.
- WHERE clause was used for specifying time frames.

Example : Recover genres and count number of titles associated with them within a specific time frame

```
SELECT DISTINCT g.idgenres, g.genre, COUNT(DISTINCT m.idmovies) AS number
FROM genres g JOIN movies_genres mg ON mg.idgenres=g.idgenres
JOIN movies m ON mg.idmovies = m.idmovies
WHERE m.year >= <startyear> AND m.year <= <endyear>
GROUP BY g.idgenres ORDER BY <sortby>
```

PostgreSQL - Query design 4/5

- Restricting to movies in case of partial matches was realized by :
 1. Looking for a number of results returned by querying the DB with a specific title
 2. Based on that information, deciding whether to access returned ID (if only a single row was returned) or performing more detailed scan (if more results were obtained) to get all titles consisted of the entered string (i.e. by applying % operators).

In case of “The Sopranos” only a SINGLE row was returned, however the exact phrase “Star Wars” is associated with more than one position in DB

PostgreSQL - Query design 5/5

Pseudocode (how it looks in practice):

```
JDBC.PerformQuery("SELECT COUNT(DISTINCT m.idmovies) AS number FROM  
movies m WHERE m.title ILIKE " + title + ";;");
```

```
if (JDBC.getResultSet().getInt("number") == 1)
```

```
JDBC.PerformQuery("SELECT DISTINCT m.idmovies FROM movies m WHERE m.title  
ILIKE " + title + ";;");
```

else

```
JDBC.PerformQuery("SELECT m.idmovies FROM movies m WHERE m.title ILIKE '%" + title + "%' AND m.idmovies NOT IN (SELECT DISTINCT m.idmovies FROM movies m  
JOIN series s ON m.idmovies = s.idmovies WHERE m.title ILIKE '%" + title + "%') AND  
m.year >= " + syear + " AND m.year <= " + eyear + " ORDER BY m." + sort + ";;");
```


PostgreSQL - Advantages

- Flexibility in acquiring any information from the DB (with a proper statement)
- Intuitive SQL statements (to an extent)

PostgreSQL - Limitations/issues

- Impedance mismatch = usually it's difficult to map java objects with relations and tables.
- More complex SQL statements can be tricky to create and “difficult to tune”.
- For related (processable) information multiple JOIN queries need to be submitted.

PostgreSQL - Lessons learned

- Linking together OOP language and RD schema can sometimes be difficult and may require a large amount of testing to prepare working solution.
- SQL statements with JOINS and their output can be messy (in terms of the structure).
- For most queries with JOINS it's good to use DISTINCT clause which usually affects its performance.

MongoDB - Schema

- Two possible approaches:
 - Copy PostgreSQL schema and perform JOINS client side
 - Denormalization
- We chose a hybrid form
- Limit denormalization to small data like keywords and genres
- Duplicating movie and actor details too inefficient

```
{  
  "_id" : ObjectId("574cb18d3233b8399595ff3d"),  
  "idmovies" : 213978,  
  "title" : "Matrix",  
  "year" : 1993,  
  "number" : "",  
  "type" : 3,  
  "location" : "",  
  "language" : "",  
  "genres" : ["action", "drama", "fantasy", "thriller"]  
}
```

MongoDB - Schema

- Base schema built by importing CSV export from PostgreSQL
- Post-processed with Python script to denormalize fields like genres and keywords based on query needs
- Indexes added to support aggregation queries for statistics operations

MongoDB - Query design

- General steps
 - Build search query to get base results (e.g. matching movie or actor ids)
 - Build bulk data query to get the details of all of the matched objects by id (simulate INNER JOIN)
- Statistics SC3 and SC5 can be implemented natively using MongoDB aggregate queries
- Add appropriate indexes to accommodate queries afterwards

MongoDB - Query design

- Example: find movies that a certain actor has played in and list their co-stars
- Steps:
 - Search query to find movie ids in the acted_in collection
 - Bulk query to find details in the movies collection using \$in operator and list of ids from previous query
 - Bulk query to find ids of actors starring in those movies using \$in query
 - Bulk query to retrieve actor details (name, gender) using all unique actor ids from previous query, again using \$in
 - Combine results client side in Java
- Fixing this would require denormalizing acted_in = massive storage requirements

MongoDB - Query design

- Client side joins through bulk queries required for nearly endpoint
- Only statistics queries (SC3 & SC5) can rely on a single query
- Required a lot of performance tuning and optimization of server - > database communication
- Queries could not be accommodated by more schema changes without significant storage increase

MongoDB - Advantages

- Document model maps nicely to Java objects and arrays
- Scaling is better supported than with PostgreSQL
- May be more suited to variable amounts of data available about movies and actors

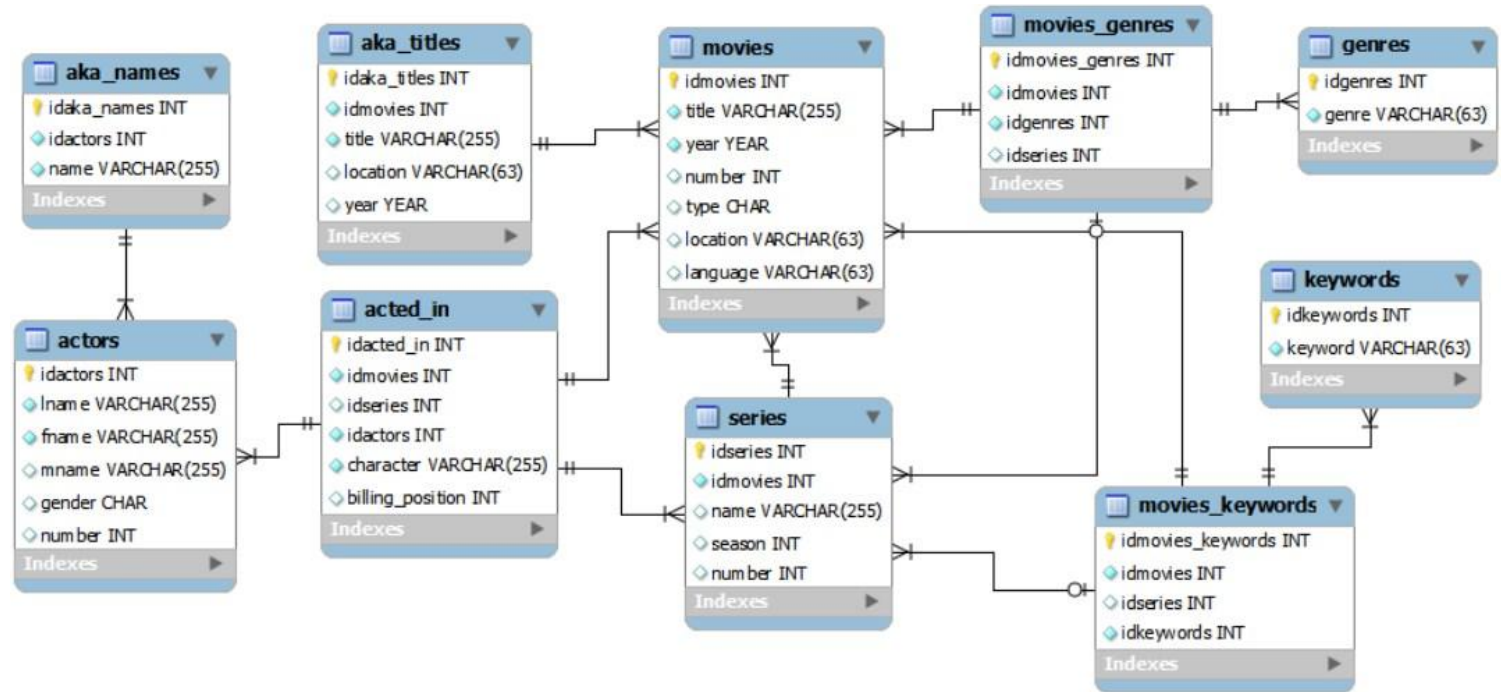
MongoDB - Limitations/issues

- PostgreSQL offers better support for hybrid approach through JSON columns
- Loose typing and optional fields in documents requires client side validation
- Impedance mismatch: MongoDB is optimized for heavy writes, not flexible reads as required by this application
- Relational database with replication much more fitting
- Increased storage requirements because of metadata

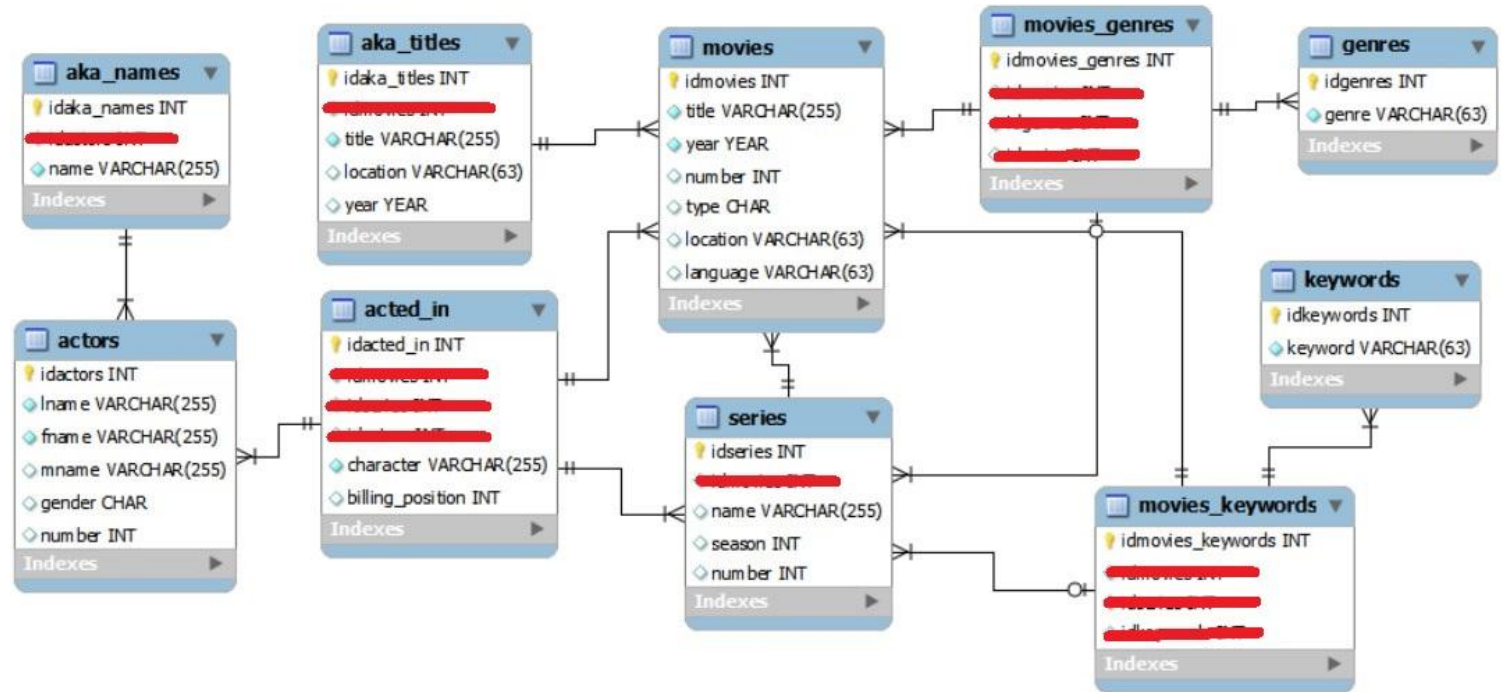
MongoDB - Lessons learned

- Less powerful queries impose constraints on data format
- Requires client side code to handle data validation and joins
- Shines when data can be denormalized well, but this is not one of those applications

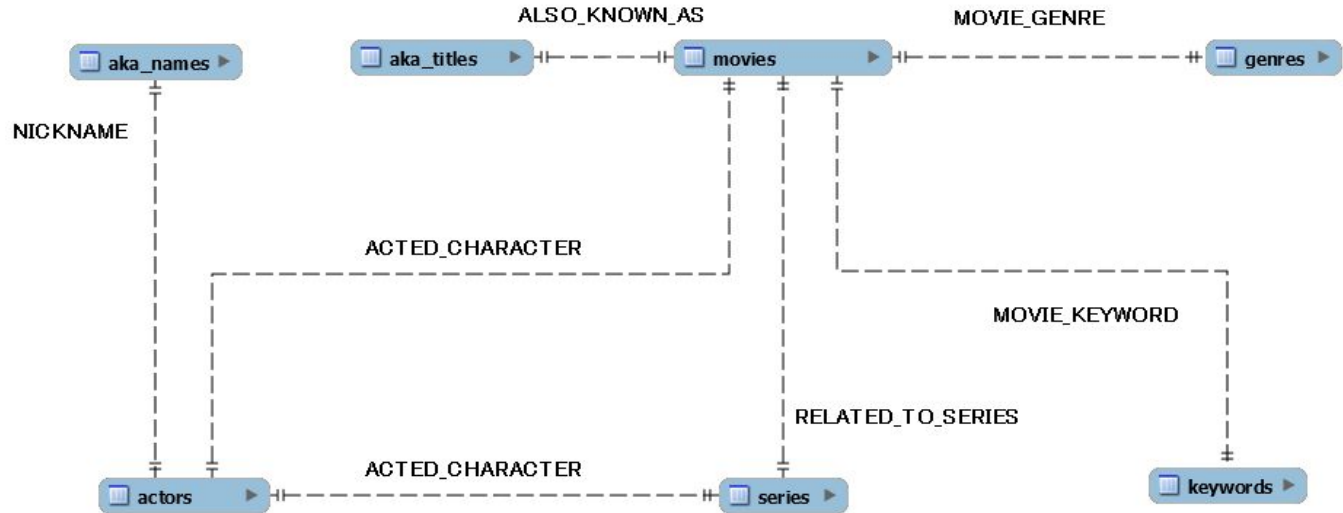
Neo4j - Schema



Neo4j - Schema



Neo4j - Schema



Neo4j - Database import

Step 1 - Extracting PostgreSQL data

```
COPY acted_in TO 'C:\acted_in.csv' WITH (FORMAT CSV, HEADER);  
COPY actors TO 'C:\actors.csv' WITH (FORMAT CSV, HEADER);  
COPY aka_names TO 'C:\aka_names.csv' WITH (FORMAT CSV, HEADER);  
COPY ...
```

Neo4j - Database import

Step 2 - Importing data

```
USING PERIODIC COMMIT 1000 LOAD CSV WITH HEADERS FROM "file:
///C:\\actors.csv" AS row CREATE (:actors{
  idactors:toInt(row.idactors),
  lname:row.lname,
  fname:row.fname,
  mname:row.mname,
  gender:toInt(row.gender),
  number:toInt(row.number)
});
```


Neo4j - Database import

Step 3 - Creating indices

```
CREATE CONSTRAINT ON (t:actors) ASSERT t.idactors IS UNIQUE;  
CREATE CONSTRAINT ON (t:aka_names) ASSERT t.idaka_names IS  
UNIQUE;  
CREATE CONSTRAINT ON (t:aka_titles) ASSERT t.idaka_titles IS  
UNIQUE;  
CREATE CONSTRAINT ON (t:genres) ASSERT t.idgenres IS UNIQUE;  
CREATE CONSTRAINT ON (t:keywords) ASSERT t.idkeywords IS  
UNIQUE;  
CREATE CONSTRAINT ON (t:movies) ASSERT t.idmovies IS UNIQUE;  
CREATE CONSTRAINT ON (t:series) ASSERT t.idseries IS UNIQUE;
```

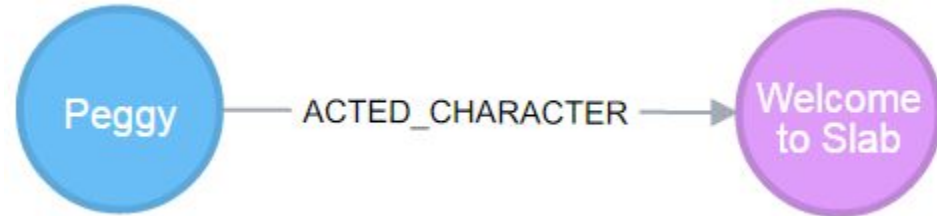
Neo4j - Database import

Step 4 - Creating relationships

```
USING PERIODIC COMMIT LOAD CSV WITH HEADERS FROM "file:///C:\acted_in.csv" AS row
MATCH (a:actors {idactors:toInt(row.idactors)})
MATCH (b:movies {idmovies:toInt(row.idmovies)})
MERGE (a)-[:ACTED_CHARACTER {
  character:coalesce(row.character, "N/A"),
  billing_position:coalesce(toInt(row.billing_position), 0)
}]->(b);
```

Neo4j

	idmovies integer	idactors integer
1	14	6
2	14	6
3	14	6
4	14	6
5	14	6
6	14	6
7	14	6
8	14	6
9	14	6
10	14	6
11	14	6
12	14	6
13	14	6
14	14	6
15	14	6
16	14	6
17	14	6
18	14	6
19	14	6



Neo4j - Query Design

- Functionally very similar to SQL
 - No need for **DISTINCT**
- Same pattern of controllers, joins replaced by relationships
 - `MATCH p=(m:movies {idmovies:1})-[r:MOVIE_GENRE]->(g:genres) RETURN p;`
- Specific info = 1 query
- Full info = up to 5 queries
 - Untapped potential?
- Sometimes, relationships are needed when joins are not

Neo4j - Advantages

- Removal of foreign keys
- Traverses relationships quickly
- Allows exotic relationships

Neo4j - Limitations/issues

Within the context of IMDB:

- Not enough relationships
- Relationships too simple

Neo4j - Limitations/issues

- Not suited for API
- Lack of parameterized queries
- Method of creating relationships (CSV) inconvenient
- Lack of different data types

Neo4j - Lessons learned

- Make sure to remove foreign keys
- Still a little immature
- Not built for supporting average use cases

Demo

- Movie information
- Actor information
- Genre information

Recap

-

Questions?

-